## CRISP 2nd annual meeting

The use of hardware virtualization in RASHPA

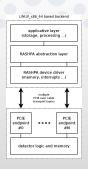


### Contents

- RASHPA context reminder
- Device virtualization
- The VPCIE framework



### Context - RASHPA reminder



RASHPA basic goal is to transfer **detector** memory contents into a **backend** for further processing, visualisation or storage. To do so, RASHPA relies on low level hardware and software components

- a data transport layer, currently PCI Express over cable
- a data transmission engine implemented on XILINX FPGA
- LINUX kernel device driver and userland software



## Context - RASHPA project issues

#### Still being prototyped

• how to simplify hardware software codesign?

Aims at using multiple PCIE links. Yet, the first prototype is single link

• how to scale the prototype platform ?

Well positionned to interface with PCIE based accelerating technologies

• how to test unavailable (too recent or expensive) hardware ?

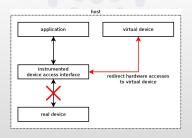


### Context - solution

Solution: device virtualization



#### Device virtualization - main idea



Applications on a **host** machine access hardware via interfaces. By **instrumenting** these interfaces, one can **redirect** the accesses to a software implementing the device. The device is said to be **virtualized**.



# Device virtualization - hardware access redirection (0)

Virtualization assumes device accesses are made using a known interface

- software library API (system calls)
- memory mapped registers
- CPU specific instructions (in, out)



# Device virtualization - hardware access redirection (1)

The interface is then instrumented to redirect access to the virtual device

- software hooking (code instrumentation, library replacing)
- instruction emulation (QEMU)
- hardware traps (page protection)
- architecture support (INTEL VT)
- paravirtualization (XEN)



## Device virtualization - applications

#### Device virtualization example applications

- support: application running on unmaintained platform
- security: sandboxing, reverse engineering
- quality: debugging, fault injection
- testing: milkymist, zinq, android



## VPCIE - goals

#### VPCIE, a Virtual PCIE framework made for RASHPA

- virtualize PCIE endpoints
- generic enough to fit other projects
- the backend software must run without any modification
- should be possible to use VHDL for device simulation



## VPCIE - opensource building blocks

#### **QEMU**

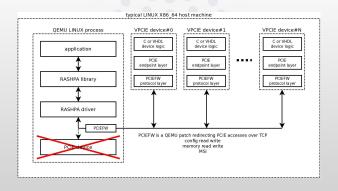
- http://wiki.qemu.org/Main\_Page
- architecture emulator (X86\_64, ARM ...)
- used to trap PCIE hardware accesses

#### **GHDL**

- http://ghdl.free.fr
- VHDL frontend for GCC
- used to implement device in VHDL

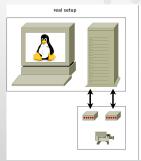


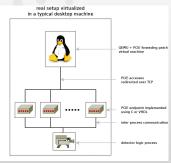
## **VPCIE** - implementation





# VPCIE - RASHPA possible setup







### VPCIE - RASHPA backend

#### RASHPA backend is a full featured LINUX system

- runs in a QEMU virtual machine
- PCIE accesses are trapped and sent over TCP to the devices
- PCIE forwarder is available as a QEMU patch



### **VPCIE - RASHPA devices**

#### RASHPA devices

- run as a LINUX processes, can be duplicated at will
- can be implemented in C or VHDL
- PCIE made simple, focus on device logic



### **VPCIE** - VHDL interface

```
entity endpoint is
 port
  rst: in std_ulogic:
  clk: in std_ulogic:
  req_en: out std_ulogic;
  rea_wr: out std_ulogic:
  req_bar: out std_ulogic_vector(pcie.BAR_WIDTH - 1 downto 0);
  req_addr: out std_ulogic_vector(pcie.ADDR_WIDTH - 1 downto 0);
  reg_data: out std_ulogic_vector(pcie.DATA_WIDTH - 1 downto 0):
  rep_en: in std_ulogic;
  rep_data: in std_ulogic_vector(pcie.DATA_WIDTH - 1 downto 0):
  mwr_en: in std_ulogic;
  mwr_addr: in std_ulogic_vector(pcie.ADDR_WIDTH - 1 downto 0);
  mwr_data: in std_ulogic_vector(pcie.PAYLOAD_WIDTH - 1 downto 0);
  mwr_size: in std_ulogic_vector(pcie.SIZE_WIDTH - 1 downto 0);
  msi_en: in std_ulogic
end entity;
```



### **VPCIE** - C interface

```
/* runtime initialization */
int pcie_init_net(pcie_dev_t*, ...);
int pcie_fini(pcie_dev_t *);
int pcie_loop(pcie_dev_t *);
/* misc config byte accessors */
int pcie_set_deviceid(pcie_dev_t*, ...);
int pcie_set_vendorid(pcie_dev_t*, ...);
/* PCIE BAR access handlers */
typedef void (*pcie_readfn_t)(uint64_t, void*, size_t, void*);
typedef void (*pcie_writefn_t)(uint64_t, const void*, size_t, void*);
int pcie_set_bar(pcie_dev_t*, ..., pcie_readfn_t , pcie_writefn_t , ...);
/* host memory read write operations */
int pcie_write_host_mem(pcie_dev_t*, uint64_t, size_t*);
int pcie_read_host_mem(pcie_dev_t*, uint64_t, size_t*);
/* send an MSI */
int pcie_send_msi(pcie_dev_t*);
```



### **VPCIE** - benefits

#### Hardware software codesign

- reduce development time
- no modification on the backend software (esp. driver)
- act as a stimuli generator for VHDL simulations

#### Platform scaling

one PCIE endpoint per LINUX process

### Investigate unavailable technologies

NVM Express support for QEMU is available



### **VPCIE** - more benefits

#### Virtual machine

- resources easily changed (RAM amount ...)
- test with different CPU architectures (x86, x86\_64, ARM ...)
- small LINUX disk images reduce reboot times (few seconds)

#### Fault injection

LINUX and software response to PCIE device disconnection



## **VPCIE** - availability

https://github.com/texane/vpcie



## VPCIE - demo

demo

